

**GPU** TECHNOLOGY  
CONFERENCE

# Image Contrast Adjustment using Nvidia Performance Primitives (NPP)

Yang Song

# Outline

- NPP Introduction
- Problem Statement
- Solution and Result
- Further Reading/Resources

# What is NPP?

- A library of GPU-accelerated image, signal and video processing functions.
- Key features:
  - Arithmetic and Logical Operations
  - Statistical Operations
  - Filter Functions, etc.
- Performance
  - 5x ~ 10x than CPU-only implementation.
- Amount
  - Around 4000 in CUDA 5.0.

# Problem Statement

- Hazy source image lacks contrast.
- Adjustment algorithm (8-bit image):
  - Offset and scale image, such that darkest pixel is mapped to 0 and brightest pixel is mapped to 255

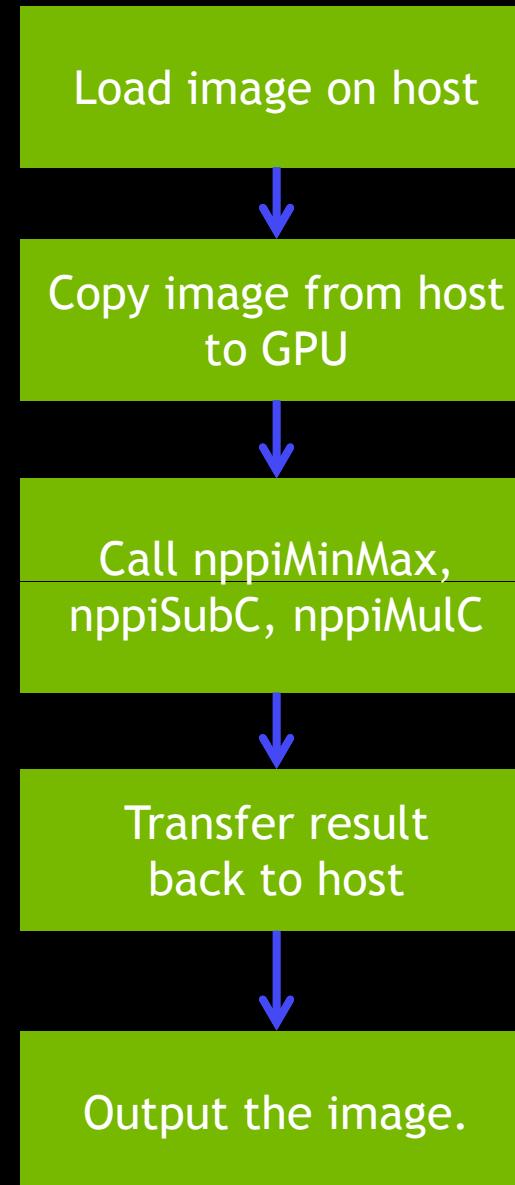
$$pDst(i,j) = 255 \times \frac{pSrc(i,j) - nMin}{nMax - nMin}$$

- Three operations needed:
  - MinMax, Subtract and Multiply



Good News!  
NPP has them all ☺

# Solution



You have to fill in  
this part.

# NPP Image Representation

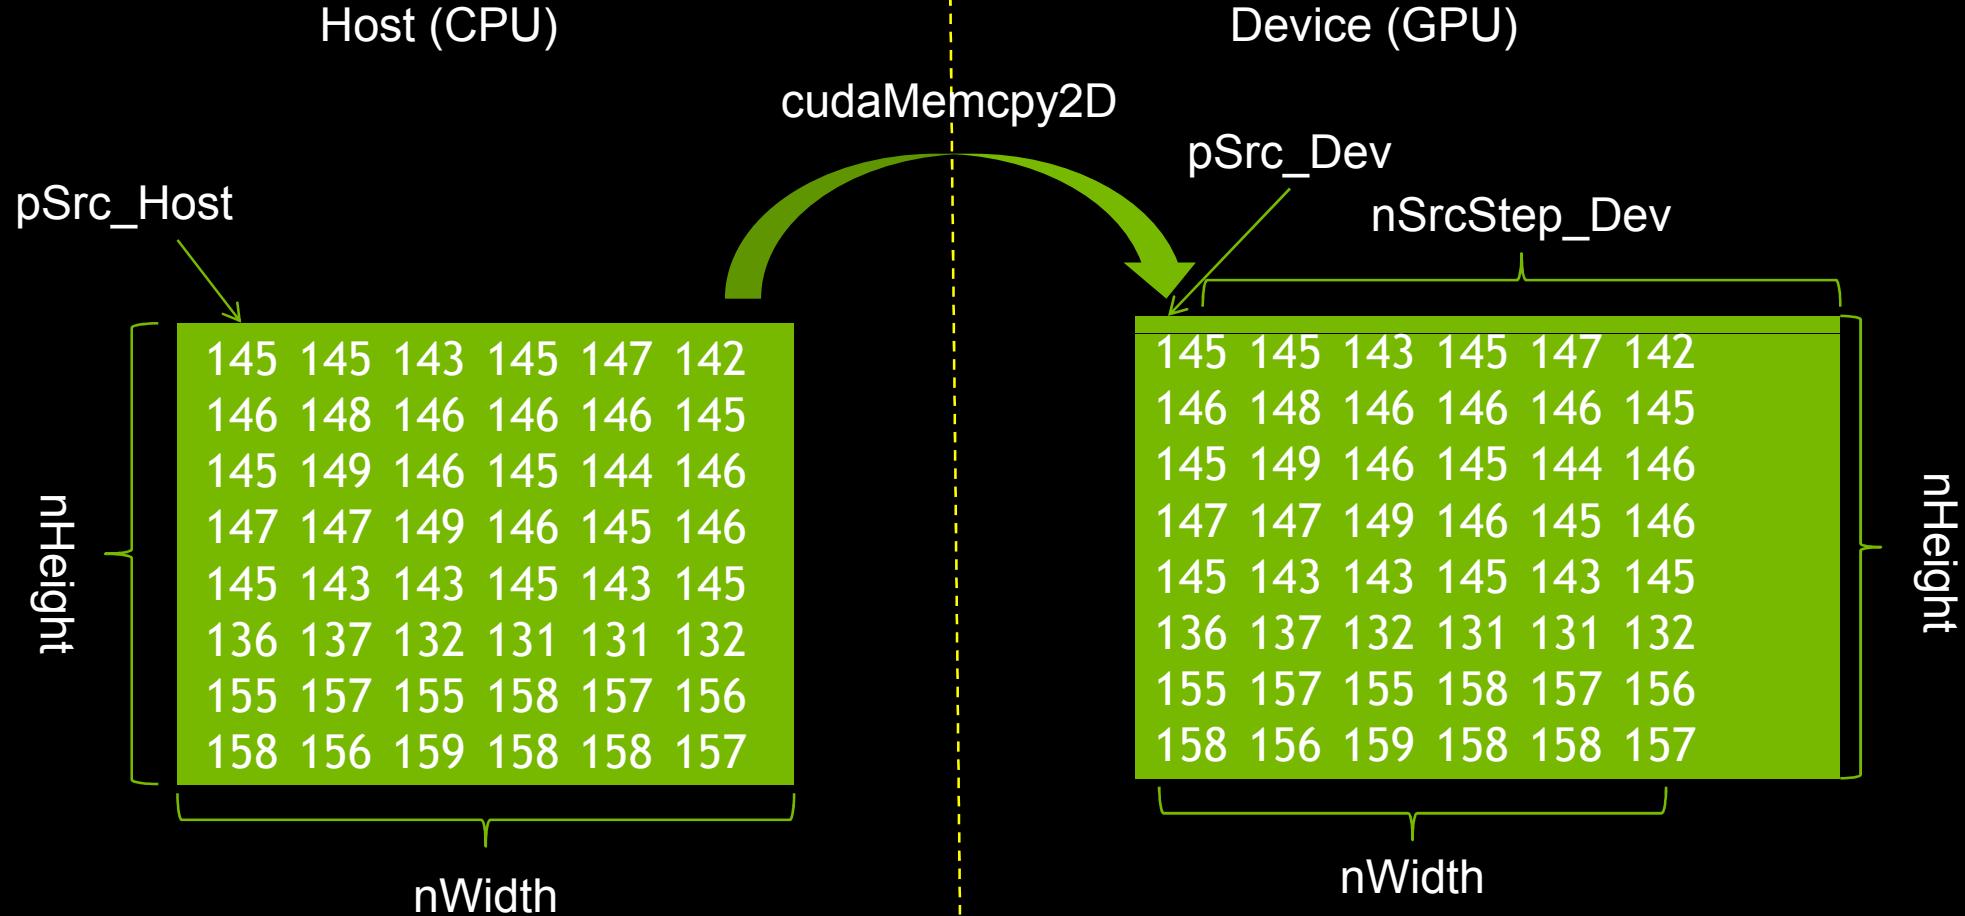
- Image is represented by two parameters:
  - pSrc: pointer to the first pixel of the image
  - nSrcStep: number of bytes between successive rows



# Function Naming

- nppiMulC\_8u\_C1IRSfs
  - npp
  - i: image module (s: signal module)
  - MulC: primitive name (Add, Sum, etc.)
  - 8u: data type of the image (16u, 32s, 32f, 64f, etc.)
  - C1: single channel (C3R, C4R, AC4R, etc.)
  - I: in-place (out-of-place if not specified)
  - R: work on ROI (region of interest)
  - Sfs: allow result scaling

# Read Image



## NPP scratch buffer

- `nppiMinMax_8u_C1R` needs device buffer
  - NPP does not allocate memory internally (unbeknownst to the user).
  - Offer users max control and flexibility on memory management.
  - Buffer can be reused to improve performance and avoid device-memory fragmentation .

## nppiMinMaxGetBufferSize\_8u\_C1R

- Parameter list:

- |            |              |                     |
|------------|--------------|---------------------|
| – NppiSize | oSizeROI     | → oROI              |
| – int *    | hpBufferSize | → &nBufferSize_Host |

## nppiMinMax\_8u\_C1R spec

- Parameter list:

- |                         |                |
|-------------------------|----------------|
| – const Npp8u * pSrc    | → pSrc_Dev     |
| – int nSrcStep          | → nSrcStep_Dev |
| – NppiSize oSizeROI     | → oROI         |
| – Npp8u * pMin          | → pMin_Dev     |
| – Npp8u * pMax          | → pMax_Dev     |
| – Npp8u * pDeviceBuffer | → pBuffer_Dev  |

## Integer-Result Scaling (Part 1)

- Avoid the clamping loss on the integer data (especially on 8u and 16u images).
- Parameter "nScaleFactor" controls the amount of scaling:  $\times 2^{-nScaleFactor}$
- Example: nppsSqr\_8u\_Sfs()
  - $255^2 = 65025$ , clamped to 255.
  - Pass scale factor 8: final result =  $255^2 \times 2^{-8} = 254.00390625$ , which will be rounded to 254.

## nppiSubC\_8u\_C1RSfs spec

- Shift minimum pixel value to 0.
  - nppiSubC\_8u\_C1RSfs
    - const Npp8u \* pSrc → pSrc\_Dev
    - int nSrcStep → nSrcStep\_Dev
    - const Npp8u nConstant → nMin, nMax, 255
    - Npp8u \* pDst → pDst\_Dev
    - int nDstStep → nDstStep\_Dev
    - NppiSize oSizeROI → OROI
    - int nScaleFactor → 0
- $$pDst(i, j) = \frac{pSrc(i, j) - nMin}{nMax - nMin} \times 255$$

# Multiplication Step

- Scale range to [0, 255].

$$pDst(i,j) = (pSrc(i,j) - nMin) \times \frac{255}{nMax - nMin}$$

- Problem: `nppiMulC_8u_C1IRSfs` uses a single Npp8u integer as the constant multiplier, but  $\frac{255}{nMax - nMin} = \frac{255}{202 - 92} = \frac{255}{110} \approx 2.318182$
- If multiply by 2 (nearest integer to actual value) error is 0.318182 (i.e. >30%!!!).
- Solution: Integer result scaling!

## Integer-Result Scaling (Part 2)

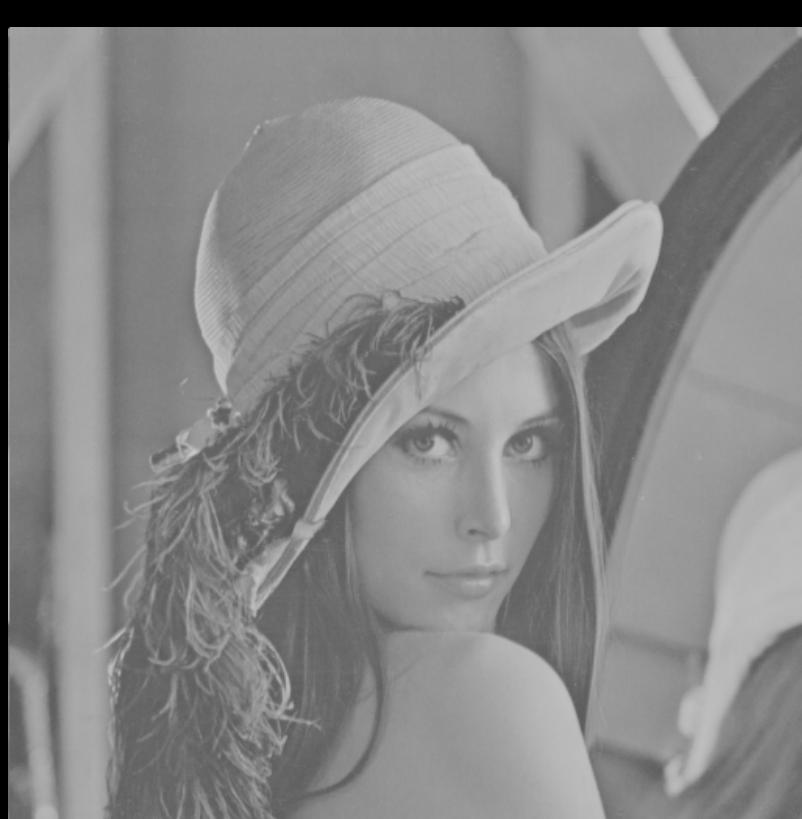
- Reduce error using Integer-Result Scaling:
  - Express exact multiplier (2.318182) as fraction of the form
$$\frac{n}{2^{-s}}$$
  - Where n is an integer and s is the integer-scale factor.
- With s=6 we get  $2^s=64$ .  $2.318182 \times 64 = 148.363648 \approx 148$
- $\frac{148}{64} = 2.3125$
- Difference with exact value  $2.318182 - 2.3125 \approx 0.005681$  which is close to the 8-bit pixel quantization error ( $1/255 \approx 0.003921$ ).

# Computing optimal nScaleFactor

```
int nScaleFactor = 0;  
int nPower = 1;  
while(nPower * 255.0f / (nMax_Host - nMin_Host) < 255.0f)  
{  
    nScaleFactor++;  
    nPower *= 2;  
}
```

## nppiMulC\_8u\_C1IRSfs spec

- nppiMulC\_8u\_C1IRSfs
  - const Npp8u nConstant → computed
  - Npp8u \* pSrcDst → pDst\_Dev
  - int nDstStep → nDstStep\_Dev
  - NppiSize oSizeROI → oROI
  - int nScaleFactor → computed



before



after

## Further Reading/Resources

- NPP is freely available as part of the CUDA Toolkit at [www.nvidia.com/getcuda](http://www.nvidia.com/getcuda).
- Source code samples demonstrating use of the NPP library:
  - Box Filter with NPP
  - Histogram Equalization with NPP
  - FreeImage and NPP Interoperability
  - Image Segmentation using Graphcuts with NPP



Thank you.  
Q&A